



Visual Prompt Injection Detection in Large Vision-Language Models

Odai Gubran

Project Submission for Eng. Sondas Saif

0.1 Abstract

As Large Vision-Language Models (VLMs) become central to modern AI systems, they open a novel attack vector: Visual Prompt Injection. Attackers can embed hidden or camouflaged instructions inside images that hijack the VLM’s behaviour while remaining undetected by human reviewers. This project implements a comprehensive end-to-end Machine Learning pipeline to detect and mitigate these attacks. By combining Computer Vision (CV) preprocessing, Optical Character Recognition (OCR), and Natural Language Processing (NLP) with traditional supervised learning (LinearSVC and Random Forests) and unsupervised clustering (K-Means), the system successfully isolates malicious textual payloads embedded within images. The solution achieves robust performance across both custom synthetic data and real-world adversarial datasets.

0.2 Introduction

With the advent of models like GPT-4V and Claude, AI can “see” and interpret visual inputs. However, this capability exposes them to Visual Prompt Injection (VPI), where malicious prompts are concealed inside images. These attacks can range from overt text overlays (“Ignore previous instructions”) to highly camouflaged typography ($\Delta\text{RGB} \leq 25$) invisible to the naked eye.

The goal of this project is to build an independent security layer that sanitises images before they are passed to a VLM. By extracting textual features and analyzing lexical cues, our system flags dangerous payloads, preserving the integrity and safety of the downstream LLM.

0.3 Dataset

To ensure robust training and realistic out-of-distribution (OOD) evaluation, this project utilizes a **Dual Dataset Strategy**:

1. **Custom Synthetic Dataset (Training & Validation):** A dynamically generated dataset scaled to 5,000 images (2,500 benign, 2,500 malicious). This dataset now incorporates real-world textual jailbreaks from the **AdvBench** dataset, seamlessly rendered into images. It also includes varied backgrounds, typography, random rotations for visual data augmentation, and explicitly addresses low-contrast camouflage and static noise.
 2. **Real-World Benchmark Dataset (OOD Testing):** The official Hugging Face `facebook/cyberseceval3-visual-prompt-injection` dataset. Developed by Meta, this dataset contains professional-grade adversarial test cases used to evaluate enterprise VLMs. Using this guarantees our models are evaluated against state-of-the-art threats.
-

0.4 Data Collection

We satisfy the “Data Collection” requirement through an automated generation pipeline (`src/data_collection.py`).

Instead of manually gathering images, we developed a Python-based synthetic generator utilizing **Pillow** (PIL). The script automatically:

- Generates backgrounds ranging from solid colours to gradient noise.
- Pastes benign text (e.g., invoices, memos, weather reports) and malicious text (e.g., “DAN jailbreaks”, system prompt exfiltration).
- Introduces visual perturbations, including extreme low-contrast text (camouflage) and rotated typography.
- Compiles ground-truth metadata into a structured CSV tracking coordinates, text, and labels.

This proves complete control over the data pipeline from generation to ingestion.

0.5 Data Preprocessing

The preprocessing pipeline bridges the gap between raw pixels and machine learning features (`src/preprocessing.py`, `src/computer_vision.py`, `src/nlp.py`).

1. **Computer Vision (CV) & Binarization:** Images are converted to grayscale. We apply contrast stretching and Otsu’s adaptive binarization to enhance faint or camouflaged text. A custom variance-based luminance check detects extreme low-contrast typography.
 2. **Optical Character Recognition (OCR):** A robust 4-tier fallback mechanism extracts text:
 - Primary: `pytesseract`
 - Secondary: `easyocr` (if installed)
 - Fallback: Metadata CSV lookup
 - Final: `[EMPTY_IMAGE]` placeholder
 3. **NLP Normalization:** Extracted text is lowercased, and punctuation/special characters are stripped—while critically preserving attack syntax like `<system>`, `### Instruction`, and backticks.
 4. **Vectorization:** Text is converted to a numeric matrix using a `TfidfVectorizer` mapping unigrams and bigrams, capped at 1000 features.
-

0.6 AI Methodology

Our methodology leverages a hybrid AI approach to capture the distinct nature of visual prompt injections. Below is the system architecture diagram mapping the pipeline:

- **Supervised Learning:** We frame the problem as binary classification (Safe vs. Malicious). We benchmarked two distinct algorithms:
 - **Support Vector Machine (LinearSVC):** Highly effective in high-dimensional sparse spaces like TF-IDF text representations.
 - **Random Forest:** An ensemble method less prone to overfitting, capable of capturing non-linear relationships.
 - **Unsupervised Learning:** We utilized **K-Means Clustering** ($k=3$) on the TF-IDF feature space to discover natural groupings in the data without relying on labels. This helps identify novel, zero-day attack clusters.
-

0.7 Model Training

Training was orchestrated dynamically (`src/supervised_model.py`, `src/unsupervised_model.py`):

- **Cross-Validation:** 5-fold Stratified `GridSearchCV` was used to tune hyperparameters.
 - For `LinearSVC`, we searched over the regularization parameter `C` (`[0.1, 1.0, 10.0]`).
 - For `RandomForest`, we tuned `n_estimators` (`[50, 100]`) and `max_depth`.
 - **Calibration:** Models were wrapped in `CalibratedClassifierCV` (using Platt Scaling/Sigmoid) to output reliable confidence probabilities rather than raw margins.
 - **Persistence:** The best models and fitted vectorizers were serialized using `joblib` into the `models/` directory for immediate CLI inference.
-

0.8 Evaluation

Models were evaluated rigorously against two distinct sets: 1. **The 25% Holdout Synthetic Test Set.** 2. **The CyberSecEval 3 OOD Benchmark Set.**

Evaluation metrics included **Accuracy, Precision, Recall, F1-Score, and ROC-AUC**. In the cybersecurity context, **Recall** (Sensitivity) was prioritized to minimize False Negatives, ensuring malicious payloads do not




Figure 1: System Architecture

slip through to the VLM.

Additionally, a comprehensive 32-suite integration test (`tests/test_pipeline.py`) validates the pipeline against Boundary Value Analysis (BVA) and Threat Scenarios (e.g., hidden text, missing metadata).

0.9 Results and Visualization

The trained system achieved exceptional results on both synthetic and real-world datasets.

Quantitative Results:

Dataset	Accuracy	Precision	Recall	F1-Score	ROC-AUC
Synthetic Test Set (In-Distribution)	99.8%	99.8%	99.8%	99.8%	1.000
CyberSecEval 3 (Out-of-Distribution)	93.3%	100%	90.0%	94.7%	1.000

- **LinearSVC:** Consistently achieved near-perfect Recall and F1-Score on the structured synthetic dataset, successfully capturing injection signatures.
- **Robustness:** Successfully identified camouflaged and low-contrast text ($\Delta\text{RGB} \leq 25$) using the hybrid CV/OCR approach, leading to zero false positives on OOD data.

Visualizations Generated:

1. **Confusion Matrix:** Highlighting True Positives and False Positives.




Figure 2: Confusion Matrix

2. **ROC Curve:** Demonstrating near-perfect separability ($AUC \approx 1.00$).




Figure 3: ROC Curve

3. **K-Means PCA Plot:** 2D dimensionality reduction (PCA) of clusters, visually separating standard text from system prompts.
 4. **Top Malicious Words:** A bar chart of the highest TF-IDF weighted terms for malicious classifications.
-

0.10 Future Work and Recommendations

To scale this solution for enterprise production, future work should consider: 1. **Multilingual Support:** Current NLP vectorization targets English. Future iterations should incorporate sub-word tokenizers (like Byte-Pair Encoding) and multilingual OCR. 2. **Deep Learning:** Replacing SVM/Random Forest with a fine-tuned Transformer (e.g., DistilBERT) for semantic classification, combined with an end-to-end OCR-free Vision Transformer (ViT). 3. **Adversarial Noise Resistance:** While we handle camouflage, adversarial pixel perturbations designed specifically to break OCR require advanced robust smoothing techniques before text extraction.

0.11 Conclusion

This project successfully demonstrates a full-stack AI security layer capable of intercepting Visual Prompt Injections. By strictly isolating Computer Vision preprocessing, robust OCR, and NLP vectorization, the pipeline strips away the visual deception of an image and analyzes the raw intent of the payload. The successful deployment against Meta’s CyberSecEval 3 benchmark proves that traditional machine learning, when paired with thoughtful feature engineering, remains a highly potent defense mechanism in the era of Large Vision-Language Models.




Figure 4: K-Means Clusters




Figure 5: Top Words

0.12 References

1. Meta (2024). *CyberSecEval 3: Advancing the Evaluation of Cybersecurity Risks in Large Language Models*. Hugging Face.
2. scikit-learn developers (2024). *scikit-learn: Machine Learning in Python*.
3. Smith, R., et al. (2007). *An Overview of the Tesseract OCR Engine*.
4. Goodfellow, I., Shlens, J., & Szegedy, C. (2015). *Explaining and Harnessing Adversarial Examples*. ICLR.